

Submitted by Howie Singer, PhD as an individual

Comments on AI Study

I am currently an Adjunct Professor in the NYU Music Business and worked for 15 years at Warner Music Group including serving as the Chief Technologist for the Digital Strategy Group. I provide consulting services on AI-related matters to both start-ups and established firms and Oxford University Press recently published a book on the history of technology disruption in the music industry that I wrote with Bill Rosenblatt entitled “Key Changes: The Ten Times Technology Transformed the Music Industry.”

Before digging into the specifics of Generative AI (Gen AI), there is an important historical precedent that is worth discussing. That is, the amount of time that it takes for the courts to resolve the thorny issues in copyright-related matters; particularly when “fair use” defenses are involved.

Disruptive technologies often outstrip the ability of legal and regulatory frameworks to keep up. Nearly 30 years elapsed after the introduction of player pianos and phonographs before the law was changed to treat piano rolls and records as copies of musical performances.

Almost a century later, Napster’s fair use defense was shot down only 20 months after the launch of the peer-to-peer file-sharing software. But it took another four years for the Supreme Court to rule against Grokster -one of Napster’s many successors- by adopting the principle of inducement from patent law and finding that they had distributed software with the intent of “promoting its users to infringe on copyright.”

At that point, most other p2p services shut down or attempted to obtain licenses from the labels to go “legit.” Limewire did not. They picked up millions more users who preferred “sharing” to paying and thrived until a federal court judge issued an injunction against them in 2010 for inducing their users to commit infringements. From start to finish, the legal battles over peer-to-peer file-sharing lasted over a decade.

The length of these battles might have been good for the law firms hired to argue on behalf of both sides, but the lack of clarity came at a cost. Paying for all those lawyers diverted money from investments in improving features at the start-ups and in signing artists at the labels. Clarity on the rules of the road for GenAI when it comes to control, consent, and compensation will accelerate its adoption by all parties.

Although there already several lawsuits over the use of artworks and written material as training material for GenAI models, there have not been any copyright lawsuits over AI models duplicating a famous person’s voice. That’s because there is no federal copyright law that covers an individual’s speaking or singing style. Some states have legal protections for name and likeness that would include vocal style. Bette Midler won a suit against Ford Motor Company in 1988 when they hired another singer to imitate her after she refused Ford’s request to use one of her performances in a commercial. The Universal Music Group’s General Counsel already stated at a Senate Judiciary Committee hearing that, to protect artists’ voices, Congress should enact a new “federal right of publicity.” It is hard to argue that anyone should be able to train and to use a Machine Learning model to duplicate someone’s voice and then create new acting or singing performances without the consent of the original artist.

In addition, rightsholders claim that the use of their creative works to train a model to generate a story, a book or lyrics without permission infringes on their copyrights. The companies building the models say they are training their models on openly available information or proprietary datasets and that any infringements, if they even exist, would be permissible under “fair use.”

There is no clear guidance to determine that any action is fair use or not. Once a party is sued for infringement, they can invoke fair use as a defense and a federal judge, based on the facts of that specific case, decides whether that argument is valid. Judges examine four factors when rendering a decision on fair use:

- The purpose and character of the use;
- The nature of the copyrighted work;
- The amount and substantiality of the portion taken;
- The effect of the use on the potential market.

Rather than waiting years for these arguments to wend their way through the courts, it would be preferable to clarify now that using copyright material for training cannot be considered fair use.

The GenAI companies believe their applications are “transformative in nature” because the resulting works neither contain pieces of the original nor are they similar enough to claim infringement. This position does not appear to be correct. For example, Anthropic can recreate almost identical copies of copyrighted lyrics. Even entering a descriptive prompt for a song about Buddy Holly’s death outputs lyrics resembling Don McLean’s “American Pie.” Under these circumstances, it is hard to deny that copyrighted material was included in the training set without permission.

Second, the tech companies will claim “fair use” by comparing how humans “learn” from the art of others. Suppose a would-be author read all of Stephen King’s works before writing her own horror novel or a budding songwriter played every Taylor Swift track on repeat before composing their new song. As long as they did not copy and re-use specific elements of the originals, this behavior is legitimate. The GenAI companies simply ask how is it different when an algorithm does exactly the same thing as those aspiring artists? Of course, there is an inherent difference in scale. GenAI incorporates more attributes than a human could remember to create works in numbers that dwarf an individual person’s abilities.

But let’s not focus merely on the sheer number of such AI creations but also on that final fair use factor: the effect they would have on the potential market for the original artists. The majority of recorded music revenues are derived from streaming and the share of plays for each track is the key factor in allocating those dollars. GenAI trained on the Bad Bunny catalog will produce facsimiles with ever-improving quality. Streaming services analyze all the tracks added to their catalog to make better recommendations for fans and playlists. The streaming services’ own tools will likely draw more attention to the AI-generated tracks based on similarities to the hits. It would be surprising if some of those did NOT attract the attention of Bad Bunny fans thus reducing the revenue share for the originals. One could say that the entities offering these voice and lyrics models are “inducing” users to create tracks specifically geared to undermine the market for the originals.

Finally, the tech companies will rely upon a rationale that has proved useful in pushing back on previous regulatory efforts. Namely, that any rules run the risk of bringing all innovation to a screeching halt. Google has already told the courts that training their chatbot Bard requires the use of public data. They

upped the ante even further by claiming that finding for the plaintiffs would "take a sledgehammer not just to Google's services but to the very idea of generative AI." This assessment is belied by history. The shutdown of p2p file sharing provided the opportunity for legitimately licensed and innovative services like Spotify to merge.

Negotiations underway for licensed uses of AI that address rightsholders concerns over control, credit, and compensation would accelerate if there were more clarity on these issues. Such agreements would encourage even broader use of AI tools by creators and fans. Investments would flow more readily if the legal risks were eliminated. The creative industries and the technology companies would all adjust to a "new normal" with GenAI as an accepted part of their businesses under a mutually acceptable set of rules.